

Vaulin_Pr1_Olympics

November 13, 2022

1 120 years of Olympic history: athletes and results

1.0.1 BI statistics course Project 1

by Nikita Vaulin, Skoltech Nikita.Vaulin@skoltech.ru

The dataset is adapted by course instructor from [kaggle](#).

```
[1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno
from statsmodels.stats.multitest import multipletests
```

Also Jupyter Notebook like to send us so many warnings such as FutureWarning and so on. Let's make our output more pretty:

```
[2]: import warnings
warnings.filterwarnings('ignore')
```

1.0.2 Reading the data

We will define here special function which uploads all the files from some with a given extension to a combined pandas dataframe

```
[3]: sep_readers = {
    ".csv": pd.read_csv,
    ".tsv": pd.read_table,
    ".hdf": pd.read_hdf,
    ".h5": pd.read_hdf,
    ".html": pd.read_html,
    ".htm": pd.read_html,
    ".xml": pd.read_xml,
    ".json": pd.read_json,
    ".xls": pd.read_excel,
    ".xlsx": pd.read_excel,
    ".xltx": pd.read_excel,
    ".ods": pd.read_excel
}
```

```

}

def read_one_table(folder, filename, extension):
    """
    Extra function to read one specific file with function suitable for a given
    ↪ file extension
    """
    file = os.path.join(folder, filename)
    return sep_readers[extension](file)

def file_reader(folder: str = 'athlete_events', ext: str = '.csv') -> pd.
    ↪ DataFrame:
    """
    This function uploads from disk a dataset divided into several files as a
    ↪ single pandas data frame.

    Several input files should contain same columns! The number of rows doesn't
    ↪ matter, the function provides
    row-wise concatenation of input files.

    Parameters:
        folder (str): name of a folder with input datasets to read
        ext (str): input files extension

    Output:
        Pandas Data.Frame combined from input files row-by-row
    """
    if not ext.startswith('.'):
        ext = f'.{ext}'
    files = [read_one_table(folder, f, ext.lower()) for f in os.listdir(folder)
    ↪ if f.endswith(ext)]
    return pd.concat(files)

```

So let's load our data. We don't need any parameters here due to their default values definition

```
[4]: df = file_reader(os.path.join('.', 'data', 'athlete_events'), '.csv')
```

Now let's proceed some initial EDA to check weather it everething alright with our data:

```
[5]: df.head(3)
```

```
[5]:
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	\
0	1	A Dijiang	M	24.0	180.0	80.0	China	CHN	
1	2	A Lamusi	M	23.0	170.0	60.0	China	CHN	
2	3	Gunnar Nielsen Aaby	M	24.0	NaN	NaN	Denmark	DEN	

	Games	Year	Season	City	Sport	\
0	1992 Summer	1992.0	Summer	Barcelona	Basketball	
1	2012 Summer	2012.0	Summer	London	Judo	
2	1920 Summer	1920.0	Summer	Antwerpen	Football	

	Event	Medal
0	Basketball Men's Basketball	NaN
1	Judo Men's Extra-Lightweight	NaN
2	Football Men's Football	NaN

Let's highlight here the numerical and categorical features.

Regarding the Year, this is neither a numerical nor a qualitative variable. Let's consider it a separate type - Time Series.

```
[6]: num_features = ['Age', 'Height', 'Weight']
cat_features = ['Name', 'Sex', 'Team', 'NOC', 'Games', 'Season', 'City', 'Sport', 'Event', 'Medal']
```

```
[7]: print(f'Here we have {len(df)} entries total in our dataframe\n')
print(f'Some summary about the data:\n')
print(df.info())
```

Here we have 271115 entries total in our dataframe

Some summary about the data:

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 271115 entries, 0 to 22390
Data columns (total 15 columns):
#   Column      Non-Null Count  Dtype
---  -
0   ID           271115 non-null  int64
1   Name         271114 non-null  object
2   Sex          271113 non-null  object
3   Age          261639 non-null  float64
4   Height       210943 non-null  float64
5   Weight       208239 non-null  float64
6   Team         271112 non-null  object
7   NOC          271111 non-null  object
8   Games        271110 non-null  object
9   Year         271108 non-null  float64
10  Season       271108 non-null  object
11  City         271108 non-null  object
12  Sport        271108 non-null  object
13  Event        271107 non-null  object
14  Medal        39782 non-null   object
dtypes: float64(4), int64(1), object(10)
```

memory usage: 33.1+ MB
None

This result also could be obtained with `df.isna().sum()` or `df.isnull().sum()` functions

According to the table all the columns expected to be numeric are actually numeric, so i don't see any numbers written as strings

1.0.3 Finding rubbish in data

We will deal with missing values later, here about the rubbish

```
[8]: for f in ['Sex', 'Season', 'Medal', 'Games', 'City', 'Sport']:  
      print(f'Unique {f} values: {df[f].unique()}')
```

```
Unique Sex values: ['M' 'F' 'G' nan]  
Unique Season values: ['Summer' 'Winter' nan]  
Unique Medal values: [nan 'Gold' 'Bronze' 'Silver']  
Unique Games values: ['1992 Summer' '2012 Summer' '1920 Summer' '1900 Summer'  
'1988 Winter'  
'1992 Winter' '1994 Winter' '1932 Summer' '2002 Winter' '1952 Summer'  
'1980 Winter' '2000 Summer' '1996 Summer' '1912 Summer' '1924 Summer'  
'2014 Winter' '1948 Summer' '1998 Winter' '2006 Winter' '2008 Summer'  
'2016 Summer' '2004 Summer' '1960 Winter' '1964 Winter' '1984 Winter'  
'1984 Summer' '1968 Summer' '1972 Summer' '1988 Summer' '1936 Summer'  
'1952 Winter' '1956 Winter' '1956 Summer' '1960 Summer' '1928 Summer'  
'1976 Summer' '1980 Summer' '1964 Summer' '2010 Winter' '1968 Winter'  
'1906 Summer' '1972 Winter' '1976 Winter' '1924 Winter' '1904 Summer'  
'1928 Winter' '1908 Summer' '1948 Winter' '1932 Winter' '1936 Winter'  
'1896 Summer' '2000 Su' '2004 Summe' nan]  
Unique City values: ['Barcelona' 'London' 'Antwerpen' 'Paris' 'Calgary'  
'Albertville'  
'Lillehammer' 'Los Angeles' 'Salt Lake City' 'Helsinki' 'Lake Placid'  
'Sydney' 'Atlanta' 'Stockholm' 'Sochi' 'Nagano' 'Torino' 'Beijing'  
'Rio de Janeiro' 'Athina' 'Squaw Valley' 'Innsbruck' 'Sarajevo'  
'Mexico City' 'Munich' 'Seoul' 'Berlin' 'Oslo' 'Cortina d'Ampezzo'  
'Melbourne' 'Roma' 'Amsterdam' 'Montreal' 'Moskva' 'Tokyo' 'Vancouver'  
'Grenoble' 'Sapporo' 'Chamonix' 'St. Louis' 'Sankt Moritz'  
'Garmisch-Partenkirchen' nan]  
Unique Sport values: ['Basketball' 'Judo' 'Football' 'Tug-Of-War' 'Speed  
Skating'  
'Cross Country Skiing' 'Athletics' 'Ice Hockey' 'Swimming' 'Badminton'  
'Sailing' 'Biathlon' 'Gymnastics' 'Art Competitions' 'Alpine Skiing'  
'Handball' 'Weightlifting' 'Wrestling' 'Luge' 'Water Polo' 'Hockey'  
'Rowing' 'Bobsleigh' 'Fencing' 'Equestrianism' 'Shooting' 'Boxing'  
'Taekwondo' 'Cycling' 'Diving' 'Canoeing' 'Tennis' 'Modern Pentathlon'  
'Figure Skating' 'Golf' 'Softball' 'Archery' 'Volleyball'  
'Synchronized Swimming' 'Table Tennis' 'Nordic Combined' 'Baseball']
```

```
'Rhythmic Gymnastics' 'Freestyle Skiing' 'Rugby Sevens' 'Trampolining'
'Beach Volleyball' 'Triathlon' 'Ski Jumping' 'Curling' 'Snowboarding'
'Rugby' 'Short Track Speed Skating' 'Skeleton' 'Lacrosse' 'Polo'
'Cricket' 'Racquets' 'Motorboating' 'Military Ski Patrol' 'Croquet'
'Jeu De Paume' nan 'Roque' 'Alpinism' 'Basque Pelota' 'Footba'
'Aeronautics']
```

I see some rubbish here:

- 'G' in Sex
- '2000 Su' and '2004 Summe' in Games
- 'Moskva' in City
- 'Footba' in Sport

For the Age:

```
[9]: print(sorted(df.Age.dropna().unique()))
mav = df.Age.dropna().sort_values().value_counts().tail(1)
print(f'\nThe most Age value is {mav.index.values.astype(int)} and it appears_
↳{mav.values} times')
```

```
[10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0, 17.0, 18.0, 19.0, 20.0, 21.0, 22.0,
23.0, 24.0, 25.0, 26.0, 27.0, 28.0, 29.0, 30.0, 31.0, 32.0, 33.0, 34.0, 35.0,
36.0, 37.0, 38.0, 39.0, 40.0, 41.0, 42.0, 43.0, 44.0, 45.0, 46.0, 47.0, 48.0,
49.0, 50.0, 51.0, 52.0, 53.0, 54.0, 55.0, 56.0, 57.0, 58.0, 59.0, 60.0, 61.0,
62.0, 63.0, 64.0, 65.0, 66.0, 67.0, 68.0, 69.0, 70.0, 71.0, 72.0, 73.0, 74.0,
75.0, 76.0, 77.0, 80.0, 81.0, 84.0, 88.0, 96.0, 97.0, 240.0]
```

The most Age value is [240] and it appears [1] times

This is an obvious outlier - the age is 240 years. The same outliers are observable for Height and Weight:

```
[10]: mhv = df.Height.dropna().sort_values().value_counts().tail(1)
lwv = df.Weight.dropna().sort_values().value_counts().tail(1)
print(f'\nThe most Height value is {mhv.index.values.astype(int)} and it_
↳appears {mhv.values} times')
print(f'\nThe least Weight value is {lwv.index.values.astype(int)} and it_
↳appears {lwv.values} times')
```

The most Height value is [340] and it appears [1] times

The least Weight value is [7] and it appears [1] times

The NOC county 3-letter codes we will check with 'noc_regions.csv' from the [kaggle](#). I will print Teams as well in order to understad what the country is behind the NOC code.

```
[11]: noc_regions = pd.read_csv(os.path.join('.', 'data', 'noc_regions.csv'))
for region in df['NOC'].unique():
    if region not in noc_regions['NOC'].to_list():
```

```
SGP
['Singapore' 'June Climene' 'Rika II' 'Singapore-2' 'Singapore-1']
nan
[]
JP
['Japan']
```

```
[12]: true_df = pd.read_csv(os.path.join('..', 'data', 'athlete_events_ideal.csv'))
      for feature in df.columns:
          print(feature, set(df[feature].unique()) - set(true_df[feature].unique()))
```

Most of these mismatches we found on our own. Let's check what are the 'Pietro Spec' in Names and 'Gymnastics M' in Events

```
[13]: df[df.Name == 'Pietro Spec']
```

I see that we will say goodbye to Mr. Pietro Spe after clearing the missing values.

```
[14]: df[df.Event == 'Gymnastics M']
```

```
[14]:
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	\
	22601	57367 Velik Nikolov Kapsazov	M	25.0	167.0	67.0	Bulgaria	BUL	

	Games	Year	Season	City	Sport	Event	Medal
	22601	1960 Summer	1960.0 Summer	Roma	Gymnastics	Gymnastics M	NaN

```
[15]: print(true_df.query('ID == 57367 and City == "Roma"').Event)
print(df.query('ID == 57367 and City == "Roma"').Event)
```

```
113240    Gymnastics Men's Individual All-Around
113241          Gymnastics Men's Team All-Around
113242          Gymnastics Men's Floor Exercise
113243          Gymnastics Men's Horse Vault
113244          Gymnastics Men's Parallel Bars
113245          Gymnastics Men's Horizontal Bar
113246          Gymnastics Men's Rings
113247          Gymnastics Men's Pommel Horse
Name: Event, dtype: object
22599    Gymnastics Men's Individual All-Around
22600          Gymnastics Men's Team All-Around
22601          Gymnastics M
0          Gymnastics Men's Horse Vault
1          Gymnastics Men's Parallel Bars
2          Gymnastics Men's Horizontal Bar
3          Gymnastics Men's Rings
4          Gymnastics Men's Pommel Horse
Name: Event, dtype: object
```

Now we see ‘Gymnastics Men’s Floor Exercise’ was lost

I also see some entry that was spoiled totally but presents in ideal df.

```
[16]: one_id = int(true_df[true_df.Name == 'Thomas Hendrikus Maria "Thom" van Dijck'].
      ↪ID.values)
df[df.ID == one_id]
```

```
[16]:
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	Games	Year	Season	\
	22463	124516 NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

	City	Sport	Event	Medal
	22463	NaN	NaN	NaN

Let’s fix all of this stuff

For sex ‘G’ obviously they are male:

```
[17]: df[df.Sex == 'G']
```

```
[17]:
```

	ID	Name	Sex	Age	Height	Weight	Team	NOC	\
42	79609	Pavel Mike	G	22.0	182.0	79.0	Czechoslovakia	TCH	
74	79630	Anatoly Mikhaylin	G	37.0	NaN	NaN	Russia	RUS	

	Games	Year	Season	City	Sport	\
42	1972 Summer	1972.0	Summer	Munich	Handball	
74	1996 Summer	1996.0	Summer	Atlanta	Sailing	

	Event	Medal
42	Handball Men's Handball	Silver
74	Sailing Mixed Two Person Keelboat	NaN

```
[18]: df.Sex.replace('G', 'M', inplace=True)
df.Sport.replace('Footba', 'Football', inplace=True)
df.Games.replace('2000 Su', '2000 Summer', inplace=True)
df.Games.replace('2004 Summe', '2004 Summer', inplace=True)
df.NOC.replace('JP', 'JPN', inplace=True)
df.Event.replace("Gymnastics M", "Gymnastics Men's Floor Exercise",
inplace=True)
```

```
[19]: df = df[df.Age != 240]
df = df[df.Height != 340]
df = df[df.Weight != 7]
```

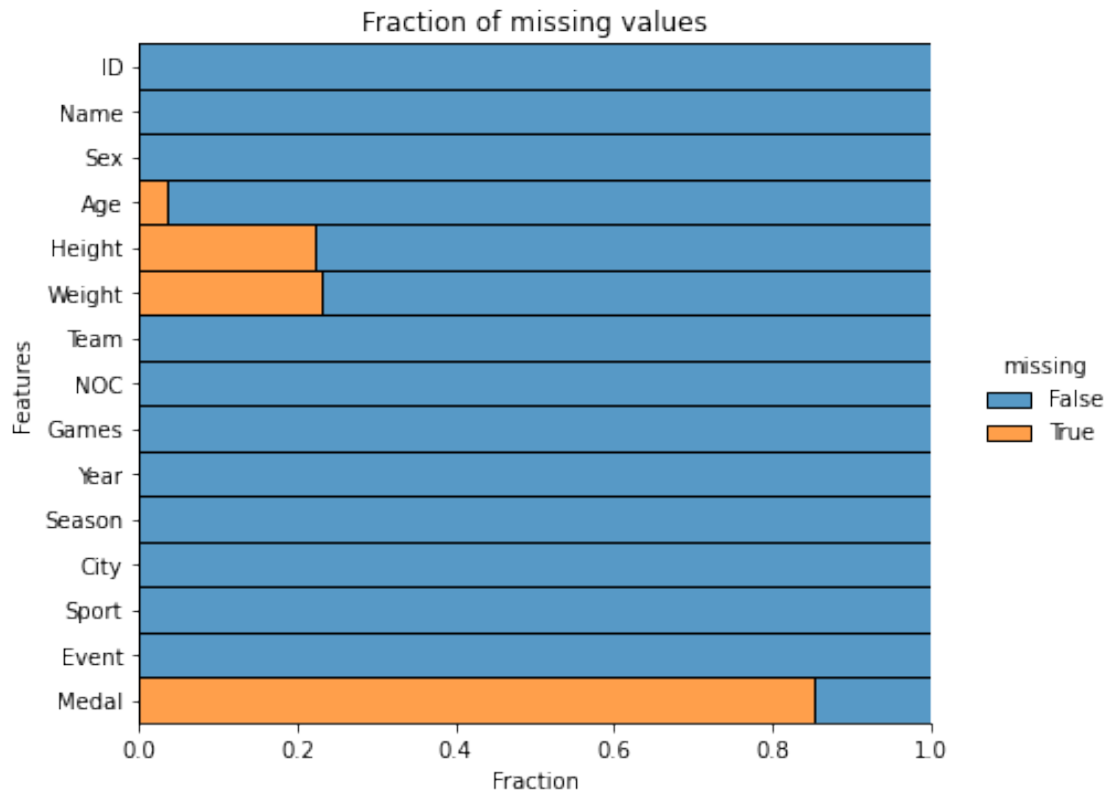
In any case, such a replacement would not be possible without an ideal dataset. In real life, you may have to restore corrupted values using the Internet or even drop them.

1.0.4 Dealing with missing values

We see there some missing values in some features. Let's make that numbers more illustratory:

```
[20]: plt.figure(figsize=(10, 6))
sns.displot(
    data=df.isna().melt(value_name="missing"),
    y="variable",
    hue="missing",
    multiple="fill",
    aspect=1.25
)
plt.xlabel("Fraction")
plt.ylabel("Features")
plt.title("Fraction of missing values")
plt.savefig("pictures/Missing_values_in_raw_dataset_count.png", dpi=100)
plt.show()
```


<Figure size 720x432 with 0 Axes>



Here we see how many missing values are there. But how are they distributed? Let's check whether there is some correlation between values' missingnesses

```
[21]: plt.figure()
fig, axes = plt.subplots(1, 2, sharex=True, figsize=(15, 5))
msno.heatmap(df, ax=axes[0], figsize=(8, 5), fontsize=12)
msno.matrix(df, ax=axes[1], sparkline=False, figsize=(6, 5), fontsize=12,
            color=(0.2, 0.10, 1.0));
plt.savefig("pictures/Missing_values_in_raw_dataset_correlation.png", dpi=100)
plt.show()
```

<Figure size 432x288 with 0 Axes>



We have some features with 1 - 2 missing values, some with 7 - 8 and 4 features with much more missing values. Firstly, let's deal with first and seconds ones.

```
[22]: df[df.Event.isna()]
```

```
[22]:
```

	ID	Name	Sex	Age	Height	Weight	Team \
22835	12005	Andrea Melissa Blackett	F	24.0	167.0	59.0	Barbados
22586	23433	Hadj Moussa Coulibaly	M	23.0	NaN	NaN	Mali
22557	34727	Carlotta Ferlito	F	17.0	160.0	50.0	Italy
22659	45919	Yuka Harada	F	28.0	170.0	60.0	Japan
22781	91137	Georgios Pantos	M	NaN	NaN	NaN	Athens-2
22409	113716	Pietro Spec	NaN	NaN	NaN	NaN	NaN
22463	124516	NaN	NaN	NaN	NaN	NaN	NaN

	NOC	Games	Year	Season	City	Sport	Event	Medal
22835	BAR	2000 Summer	NaN	NaN	NaN	NaN	NaN	NaN
22586	MLI	2004 Summer	NaN	NaN	NaN	NaN	NaN	NaN
22557	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
22659	JPN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
22781	GRE	1906 Summer	1906.0	Summer	Athina	Football	NaN	NaN
22409	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
22463	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

Here we see some obviously bad entries. There is no sense to try to fix it somehow, it is better just to drop that 8 rows on a 271115-length scale. Moreover, such features as Event, Sport, City and Season have strongest correlation in values lost so we will fix all of them at once. After that we will think about other missing values.

```
[23]: df.dropna(axis=0, subset=["Event"], inplace = True)
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 271105 entries, 0 to 22390
```

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	ID	271105 non-null	int64
1	Name	271105 non-null	object
2	Sex	271105 non-null	object
3	Age	261632 non-null	float64
4	Height	210938 non-null	float64
5	Weight	208234 non-null	float64
6	Team	271105 non-null	object
7	NOC	271105 non-null	object
8	Games	271105 non-null	object
9	Year	271105 non-null	float64
10	Season	271105 non-null	object
11	City	271105 non-null	object
12	Sport	271105 non-null	object
13	Event	271105 non-null	object
14	Medal	39782 non-null	object

dtypes: float64(4), int64(1), object(10)
memory usage: 33.1+ MB
None

Medal feature looks really intriguing with so many missing values!

```
[24]: print(df.Medal.unique())
```

```
[nan 'Gold' 'Bronze' 'Silver']
```

O-o-o! For now it's clear. 'NaN' here means no medal. Thats easy to fix.

By the way, if we had machine learning here, we would immediately turn it into a numerical categorical feature from 0 to 3

```
[25]: df.Medal.fillna("No medal", inplace = True)
print(df.Medal.unique())
```

```
['No medal' 'Gold' 'Bronze' 'Silver']
```

Let's take a clother look at Age, Height and Weight values:

For missing values in Age:

```
[26]: print(f'The {round(len(df[df.Age.isna()])*100/len(df),1)}% of enties are_
        ↳contain missing Age')
```

The 3.5% of enties are contain missing Age

Not great, not terrible, let's remove them.

```
[27]: df.dropna(subset=["Age"], axis = 0, inplace = True)
```

With height and weight everything is a little more complicated. These variables are quite important, and the number of missing values in them is too large. Let's make two

new dataframes using two different techniques, and then we'll see how it goes better. In one data frame we will completely remove all these missing values, and in the other, we will somehow replace them with.

```
[28]: print(f'The {round(len(df[df.Height.isna()])*100/len(df),1)}% of enties are_
      ↪contain missing Height')
```

The 19.7% of enties are contain missing Height

How to replace Height and Weight? Maybe we want to put here, for example, the average values. But this is impossible due to many other features in the dataset. Okay, you can say calculate the average for certain groups and substitute it depending on the group. Also incorrect - our data contains a time series. So, for time series, there are 3 main ways to fill in missing values: forward and backward filling and interpolation. In order to take into account all these moments, we will do, for example, backfilling with division into groups. As grouping variables, we will take age, gender, country and year - they should have the greatest influence on Height and Weight. Ideally, of course, we would break the age into separate intervals, but we have a lot of data and in this case we will not do feature generation and complicate things too much.

```
[29]: print(f'Mean Height is {int(df.Height.mean())}')
      print(f'Mean Weight is {int(df.Weight.mean())}')
      df_gimp = df.copy()

      df_drop = df.dropna(subset=['Height', 'Weight'])
      df_gimp.Height = df.groupby(['Age', 'Sex'])['Height'].ffill().bfill()
      df_gimp.Weight = df.groupby(['Age', 'Sex'])['Weight'].ffill().bfill()

      print(f'df_gimp shape is {df_gimp.shape}')
      print(f'df_drop shape is {df_drop.shape}')
```

Mean Height is 175

Mean Weight is 70

df_gimp shape is (261632, 15)

df_drop shape is (206158, 15)

Firts of all, let's try to work with imputed-one dataframe and the cleaned one will be reserved.

1.0.5 EDA

3. What is the age of most young sportsmens of both sexes at Olympics 1992?

```
[30]: df_gimp[df_gimp.Year == 1992].groupby("Sex")["Age"].min()
```

```
[30]: Sex
      F    12.0
      M    11.0
      Name: Age, dtype: float64
```

4. What is the mean and standart deviation of the Height for both sexes

```
[31]: df_gimp.groupby("Sex")["Height"].describe()[["mean", "std"]]
# or
df_gimp.groupby("Sex").agg({'Height': ["mean", "std"]})
```

```
[31]:
```

	Height	
	mean	std
Sex		
F	167.942479	8.794632
M	179.066542	9.322936

5. What is the mean and standard deviation of the Height for women in tennis at the Olympics 2000?

```
[32]: df_gimp.query('Year == 2000 and Sex == "F" and Sport == "Tennis"').
      .agg({'Height': ["mean", "std"]})
```

```
[32]:
```

	Height
mean	171.801587
std	6.433686

6. In what sport did the heaviest sportsmen participate at the Olympics 2006?

```
[33]: max_weight = df_gimp.query('Year == 2006').Weight.max()
      df_gimp.query('Year == 2006')[df_gimp.query('Year == 2006').Weight ==
      max_weight].Sport
```

```
[33]: 8102    Skeleton
      Name: Sport, dtype: object
```

7. How many gold medals were won by women from 1980 to 2010?

```
[34]: df_gimp.query('1980 <= Year <= 2010 and Sex == "F" and Medal == "Gold"').
      .shape[0]
```

```
[34]: 2249
```

8. How many times has John Aalberg competed in the Olympics in different years?

```
[35]: JA_all = df_gimp.query('Name == "John Aalberg"').shape[0]
      JA_games = df_gimp.query('Name == "John Aalberg"').Games.nunique()
      JA_years = df_gimp.query('Name == "John Aalberg"').Year.nunique()
      print(f'John Aalberg participated in {JA_all} competitions totally in
      {JA_games} games in {JA_years} years')
```

John Aalberg participated in 8 competitions totally in 2 games in 2 years

9. Determine the least and most represented (by number of participants) age groups of athletes at the 2008 Olympics

```
[36]: df_gimp2008 = df_gimp.query('Year == 2008')
      df_gimp2008["Age_group"] = pd.cut(x=df_gimp2008['Age'], bins=[15, 25, 35, 45,
      55], right=False)
```

```
groups_count = df_gimp2008.groupby('Age_group').Name.count()
print(f'Our age groups are: \n{groups_count}')
print(f'The least group is {groups_count.nlargest(4).index[3]} and the most is \n{groups_count.nlargest(4).index[0]}')
```

Our age groups are:

Age_group

[15, 25) 6294

[25, 35) 6367

[35, 45) 790

[45, 55) 116

Name: Name, dtype: int64

The least group is [45, 55) and the most is [25, 35)

10. How much has the number of sports changed in the 2002 Olympics compared to the 1994 Olympics?

```
[37]: def n_sports(df, year):
        return df[df.Year == year].Sport.nunique()

print(f'There are {len(set(df[df.Year == 2002].Sport) - set(df[df.Year == 1994].Sport))} sports changed')

# I suppose the sets difference is the most right way here because .nunique()
# would not take into account cases when,
# for example, no new sports appeared but there were some changes
```

There are 3 sports changed

11. Output for the winter and summer Olympics separately the top 3 countries for each type of medals

```
[38]: counts = df_gimp[['Medal', 'Season', 'NOC']].query('Medal != "No medal"').
        .groupby(['Season', "Medal"]).apply(
            lambda x: x.value_counts().head(3))
counts = counts.reset_index(level=[2, 3])[0].rename(columns={0: "Medals"})
counts
```

```
[38]:
```

Season	Medal	NOC	Medals
Summer	Bronze	USA	1185
		GER	644
		URS	596
	Gold	USA	2461
		URS	832
		GBR	616
Silver	USA	1311	
	GBR	674	

	URS	635
Winter Bronze	FIN	215
	SWE	177
	USA	161
Gold	CAN	305
	URS	250
	USA	166
Silver	USA	308
	CAN	199
	NOR	165

12. Create a new variable `Height_z_scores` and store the values of the `Height` variable into it after it has been standardized

```
[39]: def standardize(values):
        return (values - values.mean()) / values.std()

df_gimp['Height_z_scores'] = standardize(df_gimp.Height)
```

13. Create a new variable `Height_min_max_scaled` and store the values of the `Height` variable into it after it has been min-max normalized Here I could define another one my own function, but since it is an additional task I also add some to solution (the thing I also wanted do in a previous task):

```
[40]: from sklearn.preprocessing import MinMaxScaler

min_max = MinMaxScaler()
min_max.fit(df_gimp[['Height']])
df_gimp['Height_min_max_scaled'] = min_max.transform(df_gimp[['Height']])
```

14. Compare the height, weight and age of men and women who competed in the Winter Olympics

```
[41]: from scipy import stats

def honest_t_test(x, y):
    for z in [x, y]:
        if stats.shapiro(z).pvalue < 0.05:
            raise Exception('The data is not normally distributed!')
    if stats.levene(x, y).pvalue < 0.05:
        raise Exception('Samples are not homoscedantic!')
    return stats.ttest_ind(x, y)

[42]: test_1 = stats.ttest_ind(df_gimp.query('Season == "Winter" and Sex == "M").
    ↪Height,
                                df_gimp.query('Season == "Winter" and Sex == "F").
    ↪Height)
```

```
test_2 = stats.ttest_ind(df_gimp.query('Season == "Winter" and Sex == "M").
    ↪Weight,
                        df_gimp.query('Season == "Winter" and Sex == "F").
    ↪Weight)
test_3 = stats.ttest_ind(df_gimp.query('Season == "Winter" and Sex == "M").Age,
                        df_gimp.query('Season == "Winter" and Sex == "F").Age)

tests = pd.DataFrame(['Height', test_1.pvalue], ['Weight', test_2.pvalue],
    ↪['Age', test_3.pvalue]),
                    columns=['Feature', 'p_value'])
tests['p_value_adj'] = [float(f'{x :.2g}') for x in multipletests(tests.p_value.
    ↪values, alpha=0.05, method='hs')[1]]
tests['p_value'] = [float(f'{x :.2g}') for x in tests['p_value']]
tests['Significant_differences'] = tests['p_value_adj'] < 0.05
tests = tests.set_index('Feature')
```

Here we tested 3 hypotheses, conducted p-value adjustments, rounded results and labeled whether there are significant differences between males and females. Now we can print it for publishing, for example, as a markdown table

```
[43]: print(tests.to_markdown())
```

Feature	p_value	p_value_adj	Significant_differences
Height	0	0	True
Weight	0	0	True
Age	1.2e-223	1.2e-223	True

Do height, weight and age significantly differ for men and women who competed in the Winter Olympics?

Feature	p_value	p_value_adj	Significant_differences
Height	0	0	True
Weight	0	0	True
Age	1.2e-223	1.2e-223	True

15. Are the Team and the Medal related? Let's build a contingency table

```
[44]: ct = pd.crosstab(df_gimp.Team, df_gimp.Medal, margins=True)
      ct.iloc[-5:]
```


```
[44]: Medal      Bronze   Gold  No medal  Silver    All
Team
Zefyros         0        0         2        0        2
Zimbabwe         1       17       284        4       306
Zut              0        0         0        3        3
rn-2             0        0         5        0        5
```



```
All      13006  13224    222582   12820   261632
```

```
[45]: sns.heatmap(ct.iloc[0:len(ct) - 2].query('Bronze + Silver + Gold >= 1000')[['Silver', 'Bronze', 'Gold']], vmax=1000)
```

```
[45]: <AxesSubplot: xlabel='Medal', ylabel='Team'>
```




Here we see some teams have more medals than the others, but we need somehow to prove it. Additionally, of course, we need take into account number of games a team played. Thus, let's normalize each team's values.

```
[46]: ct_normalized = ct.apply(lambda x: x / ct.All, axis=0)
ct_normalized.iloc[-5:]
```

```
[46]: Medal      Bronze      Gold  No medal      Silver  All
Team
Zefyros    0.000000  0.000000  1.000000  0.000000  1.0
Zimbabwe    0.003268  0.055556  0.928105  0.013072  1.0
Zut         0.000000  0.000000  0.000000  1.000000  1.0
rn-2        0.000000  0.000000  1.000000  0.000000  1.0
All         0.049711  0.050544  0.850745  0.049000  1.0
```

```
[47]: ct_normalized[['Silver', 'Bronze', 'Gold']].plot.bar(stacked=True)
```

```
[47]: <AxesSubplot: xlabel='Team'>
```



Once again, we see medals not uniformly distributed. The teams are not clearly seen, yeah, this plot is not for publishing.

However, we need to do a statistical analysis of this. We use Pearson's chi-square test for cross tables. It could even be calculated manually, but I know a package that allows you to build contingency tables immediately with the calculation of these statistic. As before, we will normalize data by rows (only in this case, the normalization is not in fractions, but in percentages).

```
[48]: import researchpy as rp

ct_rp, ct_chi2_test = rp.crosstab(df_gimp.Team, df_gimp.Medal, prop='row',
    ↳test='chi-square')
ct_rp.head()
```

```
[48]:
```

	Medal				
Medal	Bronze	Gold	No medal	Silver	All
Team					
30. Februar	0.00	0.0	100.00	0.0	100.0
A North American Team	100.00	0.0	0.00	0.0	100.0
Acipactli	0.00	0.0	100.00	0.0	100.0
Acturus	0.00	0.0	100.00	0.0	100.0
Afghanistan	2.56	0.0	97.44	0.0	100.0

```
[49]: ct_chi2_test
```

```
[49]:
```

	Chi-square test	results
0	Pearson Chi-square (3324.0) =	41021.6856
1	p-value =	0.0000
2	Cramer's V =	0.2286

```
[50]: if ct_chi2_test.results[1] < 0.05:
    dependence = 'some'
else:
    dependence = 'no'

print(
    f'The pearson chi-square statistics is {round(ct_chi2_test.results[0], 2)}
    ↳with p-value {ct_chi2_test.results[1]}')
print(f'That means there are {dependence} statistical dependence between Teams
    ↳and Medals')
```

The pearson chi-square statistics is 41021.69 with p-value 0.0
That means there are some statistical dependence between Teams and Medals

1.0.6 16. Additional hypotheses

Athletes height through the years

```
[51]: pearr = stats.pearsonr(df_gimp.Height, df_gimp.Year)

if pearr.pvalue < 0.05:
    dependence = 'some'
else:
    dependence = 'no'

if pearr.statistic < 0:
    changes = 'decrease'
```

```

else:
    changes = 'increase'

if abs(pearr.statistic) > 0.6:
    significance = 'significantly'
else:
    significance = 'slightly'

print(f'The pearson r coefficient is {round(pearr.statistic, 3)} with p-value_{pearr.pvalue:.2g}')
print(f'That means there are {dependence} statistical dependence between_{athletes Height and Year}')
print(f'Moeover, we can say that athletes Height {significance} {changes}_{through the years}')

```

The pearson r coefficient is -0.061 with p-value 4.6e-212

That means there are some statistical dependence between athletes Height and Year

Moeover, we can say that athletes Height slightly decrease through the years

Does the first letter in the name matter?

```

[52]: df_gimp['First_letter'] = df_gimp.Name.apply(lambda x: x[0].upper())
ct_rp, ct_chi2_test = rp.crosstab(df_gimp.First_letter, df_gimp.Medal,
    prop='row', test='chi-square')

if ct_chi2_test.results[1] < 0.05:
    dependence = 'some'
else:
    dependence = 'no'

print(
    f'The pearson chi-square statistics is {round(ct_chi2_test.results[0], 2)}_{with p-value {ct_chi2_test.results[1]}}')
print(f'That means there are {dependence} statistical dependence between first_{letter in the name and medals}')

sns.heatmap(ct_rp.iloc[1:len(ct_rp) - 1, 0:4], vmax=30);
plt.ylabel("First letter in the name")
plt.xlabel("Medal")

```

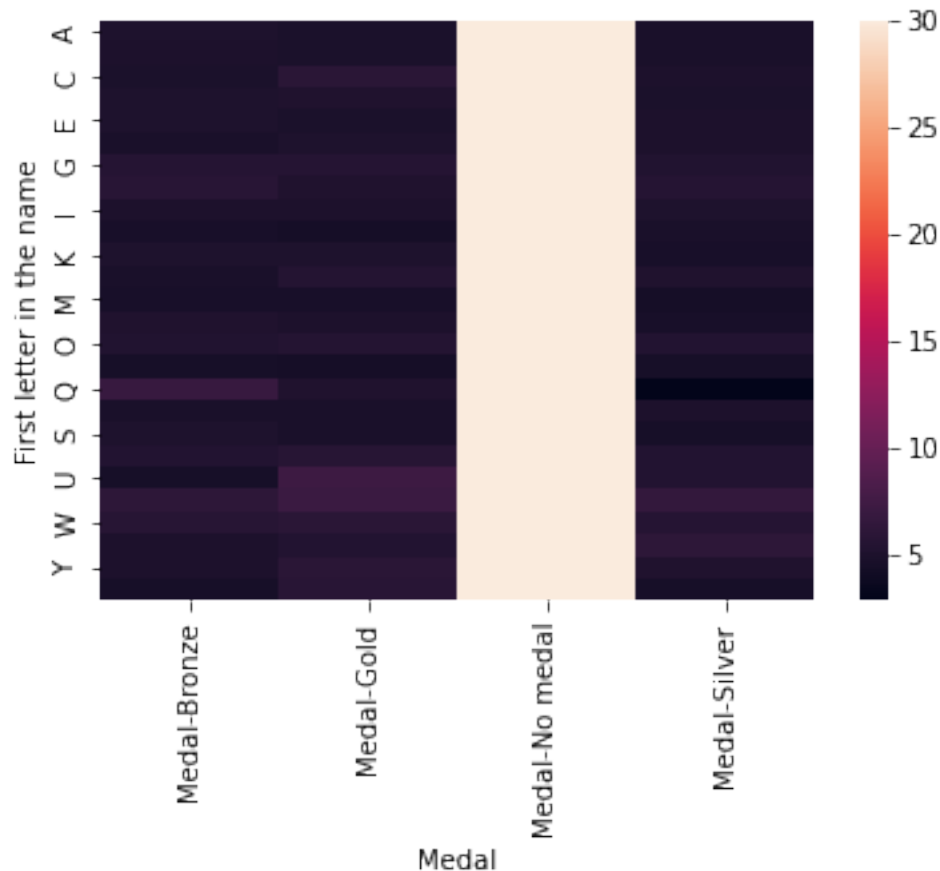
The pearson chi-square statistics is 390.01 with p-value 0.0

That means there are some statistical dependence between first letter in the name and medals

```

[52]: Text(0.5, 15.0, 'Medal')

```



Let's take a closer look at what is that - to be a Nikita

```
[53]: df_gimp['Is_Nikita'] = df_gimp.Name.str.contains('Nikita')

ct_rp, ct_chi2_test = rp.crosstab(df_gimp.Is_Nikita, df_gimp.Medal, prop='row',
    ↪test='chi-square')

if ct_chi2_test.results[1] < 0.05:
    dependence = 'some'
else:
    dependence = 'no'

print(
    f'The pearson chi-square statistics is {round(ct_chi2_test.results[0], 2)}_
    ↪with p-value {ct_chi2_test.results[1]}')
print(f'That means there are {dependence} statistical dependence between_
    ↪Nikitas and medals')


sns.heatmap(ct_rp.iloc[0:len(ct_rp) - 1, 0:4], vmax=10);
```

```
plt.xlabel("Medal")
plt.ylabel("Is Nikita?")
```

The pearson chi-square statistics is 0.23 with p-value 0.9719

That means there are no statistical dependence between Nikitas and medals

```
[53]: Text(33.0, 0.5, 'Is Nikita?')
```



As we can see, the p-value is quite large and there is not much statistical significance, but on the heatmap (despite the fact that the sizes of the groups are highly disproportional) with adjusted brightness, it can be seen that Nikita's medal strip is a little bit brighter. Noone can hide the truth!